

Course Notes for FSML Part II*

Tobías Chavarría

DSTI | DSBD2-001

Contents

Introduction	2
Chapter 1: Estimation for one parameter	2
Introduction	2
Point estimation	9
Method of moments	10
Applications	11
The Maximum Likelihood	14
Applications	15
Property	21
Theorem	23
Quality of an estimator	23
Application	24
Theorem	26
Improve estimators.	26
Exercise	27
Practical Work	28
Day 3 (In progress)	38
Confidence interval	38
Chi squared distribution	42
t-student distribution	43
Day 4 (In progress)	44
Use of t-student table	44
Confidence interval	44
Appendix	46
Types of convergence	46

*Replication files are available on the author's Github account (https://github.com/tobias-chc/FSML_II).

Introduction

Statistics notation:

1. If X_1, \dots, X_n are random variables (r.v).
2. x_1, \dots, x_n are observations.
3. If we write *i.i.d* means that the r.v are independent and identically distributed.

First aim: To propose a model for a random variable.

Generalization to multi-dimensional case:

- Y : response variable.
- $X^{(1)}, \dots, X^{(p)}$: explanatory variables.

Aim: To find a functional link between Y and the explanatory variables.

To find this functional link, the method to apply depends on the nature of the r.v's.

Y	Model
Numeric	Linear model
Qualitative (labels)	Classification

Linear model

A linear model is given by:

$$Y_i = \beta_0 + \beta_1 X_i^1 + \dots + \beta_p X_i^p + \varepsilon_i$$

where:

- β_0, \dots, β_p are unknown *fixed* parameters that can be estimated by two methods:
- Point estimation
- Confidence interval
- ε is the noise and also a random variable.

Chapter 1: Estimation for one parameter

Previous Knowledge

- Random Variable:
- The notion of distribution.
- The expectation and variance
- The distribution function
- The classical distributions (in particular the Gaussian)
- The Law of Large numbers and the Central Limit theorem

Introduction


Given x_1, \dots, x_n numeric observations, to try to find a correct parametric model, we can use 2 [graphs](#):

Plot type	Variable type	Density
Bar plot	Discrete	count/n
Histogram	Continuous	count for a bin/ n x (length of the bin)

Barplot for discrete variables

```
max.temp <- c(22, 27, 26, 24, 23, 26, 28)


barplot(max.temp,
        main = "Maximum Temperatures in a Week",
        xlab = "Degree Celsius",
        ylab = "Day",
        names.arg = c("Sun", "Mon", "Tue", "Wed", "Thu", "Fri", "Sat"))
```




Histogram for continuous variables

```
A <- rnorm(500, 0, 1)

## The default execution of this function doesn't generate a density:
hist(A)
```




```
## You need to set freq = FALSE:  
hist(A, freq = FALSE)
```




```
## You can set the numbers of bins that you want to use:  
hist(A, freq = FALSE, breaks = 100)
```

Histogram of A



*# But in order to create more breaks, you need to increase the
numbers of observations:*

```
A <- rnorm(50000, 0, 1)
hist(A, freq = FALSE, breaks = 100)
```

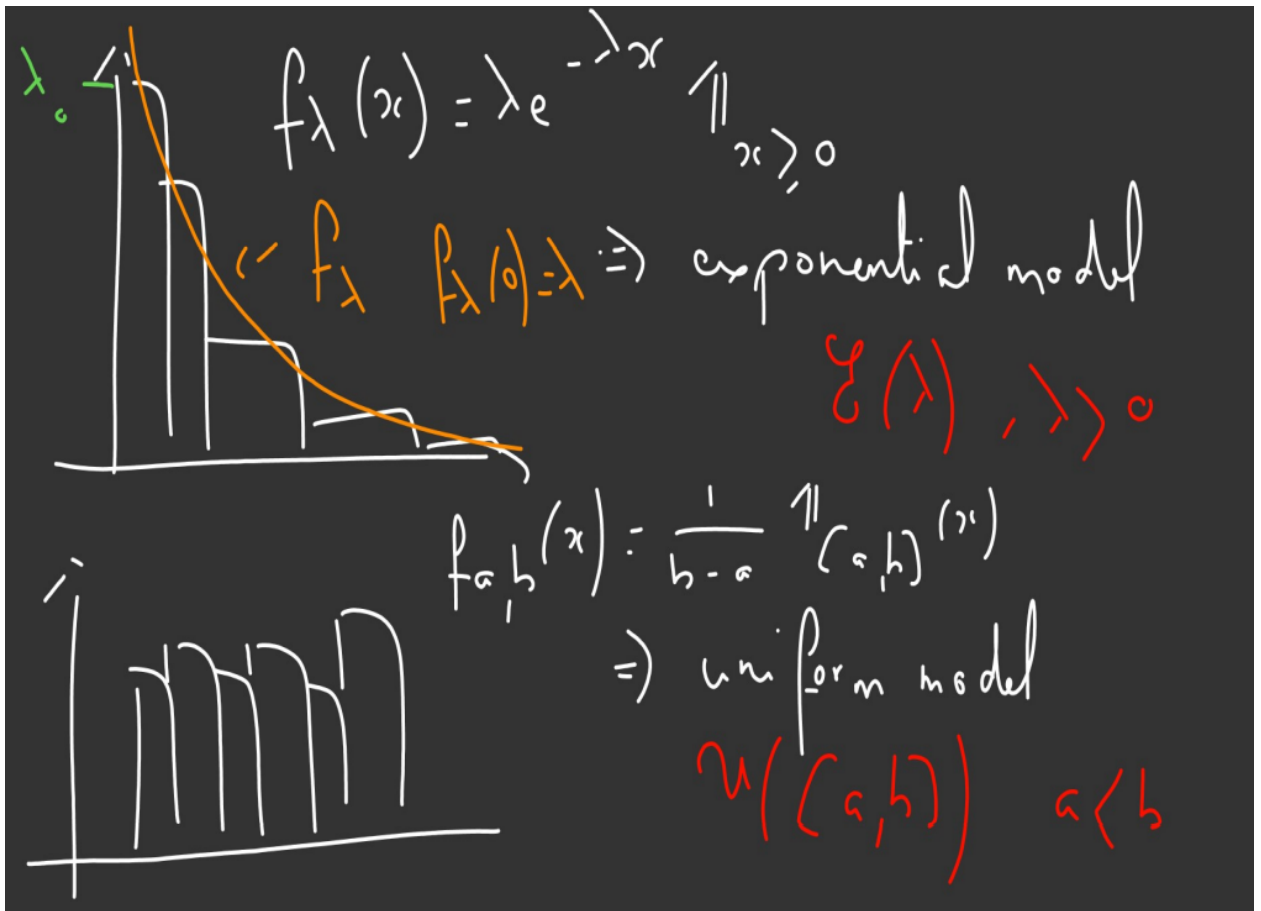


To propose a parametric model:

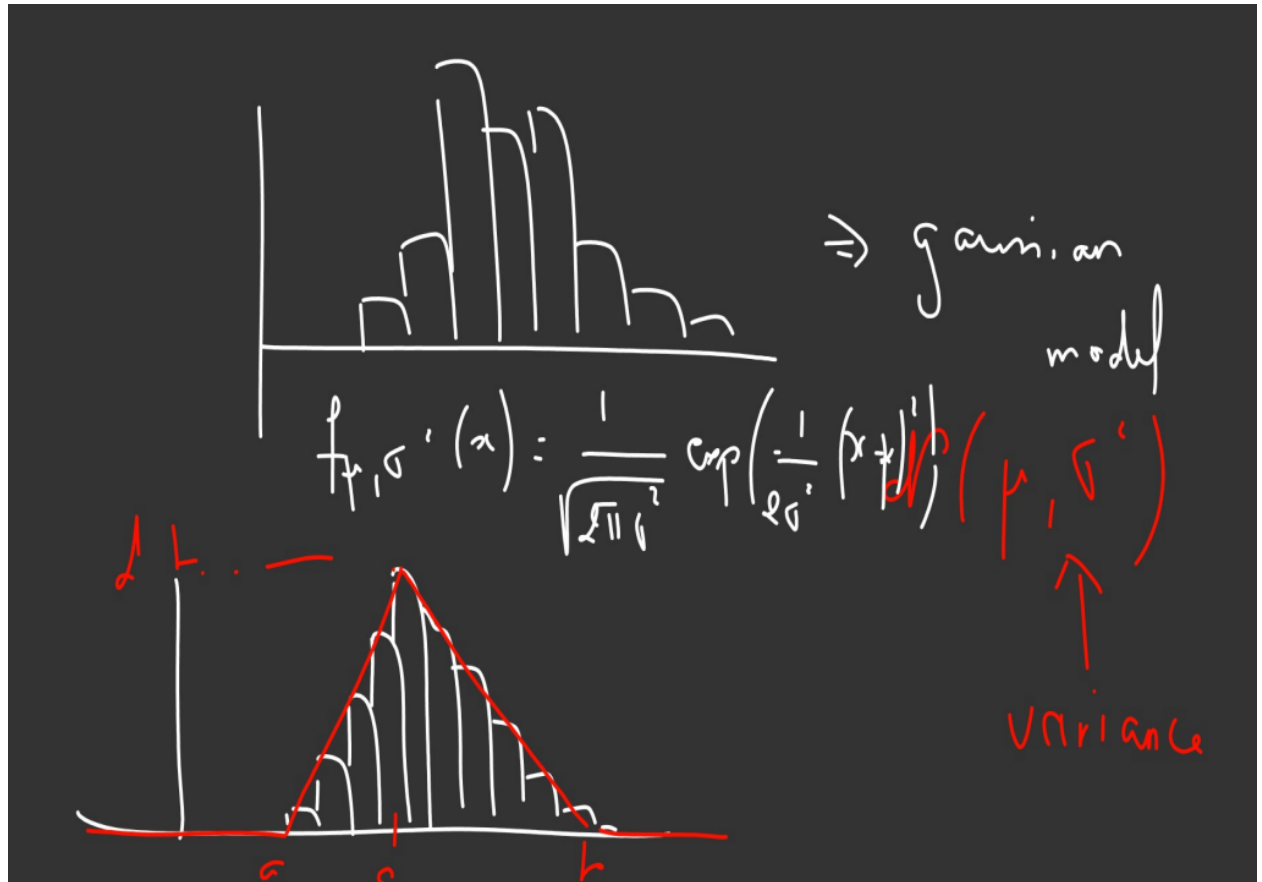
1. Make a graphical representation of the observations.
2. Guess a theoretical model by looking the previous graphic.

Examples:

Pasarlos a R!!



Example 1:



Example 2:

Question: with a representation, we can guess a parametric family of models, denoted by $\{P_\theta, \theta \in \Theta\}$. How to guess a correct value for θ thanks to the observations?

Answer: *Estimation.*

Point estimation

Let x_i an observation of a r.v X_i we assume that X_1, \dots, X_n are *i.i.d* with common distribution P_θ .

Estimator Definition: An estimator of Θ is just a function of X_1, \dots, X_n that **does not depend onto others unknown parameters.**

Remark: An estimator is a random variable!

Estimation Definition: An estimation is the value of an estimator computed thanks to the observations.

Example

Consider X_1, \dots, X_n exponential distributed and *i.i.d*, an *estimator* of λ is $\hat{\lambda}_n = \frac{n}{\sum X_i}$ an *estimation* is

$$\hat{\lambda}_n = \frac{n}{\sum x_i}.$$

Distribution	Parameter	Estimator	Estimation
Exponential $\xi(\lambda)$	λ	$\frac{n}{\sum X_i}$	$\frac{n}{\sum x_i}$

Bias (for univariate parameter) Definition: Let consider $\hat{\theta}_n$ an estimator of θ .

The bias of $\hat{\theta}_n$ is defined by:

$$b(\hat{\theta}_n) := \mathbb{E}(\hat{\theta}_n) - \theta$$

- We say that $\hat{\theta}_n$ is an unbiased estimator if

$$\forall n \in \mathbb{N}^+ \quad b(\hat{\theta}_n) = 0$$

- We say that $\hat{\theta}_n$ is *asymptotic unbiased* estimator if:

$$b(\hat{\theta}_n) \rightarrow 0 \text{ as } n \rightarrow +\infty$$

How to construct estimator?

- Method of moments
- less computations
- based on the Law of large numbers
- Maximum likelihood

Method of moments Let θ a parameter to estimated, parameter which is associate to X_1, \dots, X_n i.i.d r.v.

Let consider $k \in \mathbb{N}^*$:

- the moment of order k : $\mathbb{E}[X^k]$
- the centered moment of order k : $\mathbb{E}[X - \mathbb{E}[X]]^k$

If there exist a value k such that:

- (a) $\mathbb{E}[X^k] = g(\theta)$
- (b) $\mathbb{E}[(X - \mathbb{E}[X])^k] = h(\theta)$

Then the estimator $\hat{\theta}_n$ of θ is solution of:

- (a) $g(\hat{\theta}_n) = \frac{1}{n} \sum X_i^k$
- (b) $h(\hat{\theta}_n) = \frac{1}{n} \sum (X_i - \bar{X}_n)^k$

where

$$\bar{X}_n = \frac{1}{n} \sum X_i$$

is the **empirical mean**.

Remark: Exponential Distribution

If Let $X \sim \xi(\lambda)$, then:

- $f_\lambda(x) = \lambda \exp(-\lambda x) \mathbf{1}_{x \geq 0}$
- $\mathbb{E}[X] = \int_{\mathbb{R}} x \cdot f_\lambda(x) dx$
- $V[X] = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$

Transfer formula:

$$\mathbb{E}[l(X)] = \int_{\mathbb{R}} l(x) \cdot f_\lambda(x) dx$$

Applications

1. Let consider X_1, \dots, X_n exponential distributed and *i.i.d* compute an estimator of λ using the methods of moments.

Solution:

Let $X \sim \xi(\lambda)$, so:

$$\mathbb{E}[X] = \frac{1}{\lambda} \quad \text{and} \quad V[X] = \frac{1}{\lambda^2}$$

By applying the method of moments ($k = 1$) we get:

$$\frac{1}{\hat{\lambda}_{n,1}} = \frac{1}{n} \sum X_i$$

Thus:

$$\hat{\lambda}_{n,1} = \frac{n}{\sum X_i}$$

in the same way but using the variance ($k = 2$), we get:

$$\frac{1}{\hat{\lambda}_{n,2}^2} = \frac{1}{n} \sum (X_i - \bar{X}_n)^2$$

Thus:

$$\hat{\lambda}_{n,2} = \frac{\sqrt{n}}{\sqrt{\sum (X_i - \bar{X}_n)^2}}$$

```
A = rexp(500, 4)
```

```
1/mean(A)
```

```
## [1] 3.953248
```


```
m = c()

for (i in 1:50) {
  A = rexp(500, 4)
  m[i] <- 1/mean(A)
}

mean(m)
```

```
## [1] 3.988484
```

```
boxplot(m)
```



```
## With more observations we got less variation (500 -> 5000)
```

```
## Law of large numbers
```

```
m = c()

for (i in 1:50) {
  A = rexp(50000, 4)


```

```
m[i] <- 1/mean(A)
}
```

```
mean(m)
```

```
## [1] 4.002049
```

```
boxplot(m)
```



2. Let consider X_1, \dots, X_n *i.i.d* $\mathcal{U}([0, \sigma])$.

- Determine an estimator of σ using the methods of moments.
- Let denote $\hat{\sigma}_n = 2\bar{X}_n$. Is $\hat{\sigma}_n$ an unbiased estimator?

Solution:

Let $X \sim \mathcal{U}([0, \sigma])$.

$$\mathbb{E}(X) = \int_{\mathbb{R}} x \cdot f_X dx = \int_{\mathbb{R}} x \cdot \frac{1}{\sigma} \cdot \mathbf{1}_{[0, \sigma]} dx = \frac{1}{\sigma} \int_0^{\sigma} x dx = \frac{1}{\sigma} \frac{x^2}{2} \Big|_0^{\sigma} = \frac{\sigma}{2}$$

By the method of moments, we get that an estimator $\hat{\sigma}_n$ is solution of:

$$\frac{\hat{\sigma}_n}{2} = \frac{1}{n} \sum X_i = \bar{X}_n$$

Thus:

$$\hat{\sigma}_n = 2\bar{X}_n$$

since $\hat{\sigma}_n$ it is a function of X_1, \dots, X_n it's an estimator.

Let's compute the bias for this estimator, let $n \in \mathbb{N}^*$.

$$b(\hat{\sigma}_n) = \mathbb{E}(\hat{\sigma}_n) - \sigma$$

First, compute the expected value of $\hat{\sigma}_n$:

$$\begin{aligned} \mathbb{E}(\hat{\sigma}_n) &= \mathbb{E}(2\bar{X}_n) \\ &= \mathbb{E}\left[\frac{2}{n} \sum X_i\right] \\ &= \frac{2}{n} \sum \mathbb{E}[X_i] \quad (\text{by linearity}) \\ &= \frac{2}{n} \cdot n \cdot \mathbb{E}[X_1] \quad (\text{because are identically distributed}) \\ &= \sigma \end{aligned}$$

Therefore $b(\hat{\sigma}_n) = 0$ and then we conclude that $\hat{\sigma}_n = \bar{X}_n$ is an unbiased estimator for σ .

3. Let consider X_1, \dots, X_n *i.i.d* $\mathcal{U}([- \sigma, \sigma])$.

Since $\mathbb{E}[X_1] = 0$ (not a function of σ)

we can't use the first moment, then we use the second moment:

$$V[X_1] = \frac{\sigma^2}{3}$$

then we have:

$$\hat{\sigma}_n = \sqrt{\frac{3}{n} \sum (X_i - \bar{X}_n)^2}$$

The Maximum Likelihood

Likelihood Definition: Let X_1, \dots, X_n independent random variables, whose distributions are all depending on the same parameter θ .

Let x_1, \dots, x_n observations of those r.v

$$\mathcal{L}(x_1, \dots, x_n, \theta) = \begin{cases} \prod_{i=1}^n P_\theta(X_i = x_i) & (\text{discrete framework}) \\ \prod_{i=1}^n f_{X_i, \theta}(X_i = x_i) & (\text{continuous framework}) \end{cases}$$

where $f_{X_i, \theta}$ are the density function of X_i .

Estimator thanks to the maximum likelihood **Definition:** $\hat{\theta}_n$, an estimator for θ , due to the maximum likelihood, is solution of:

$$\mathcal{L}(x_1, \dots, x_n, \theta) = \max_{\theta} \mathcal{L}(x_1, \dots, x_n, \theta)$$

Applications

1. Let consider $X_1, \dots, X_n \xi(\lambda)$ *i.i.d.* Compute the maximum likelihood estimator.

Solution:

Let $x_1, \dots, x_n \in \mathbb{R}$.

$$\begin{aligned} \mathcal{L}(x_1, \dots, x_n, \theta) &= \prod_{i=1}^n \lambda e^{-\lambda x_i} \cdot \mathbf{1}_{x_i \geq 0} \\ &= \lambda^n e^{-\lambda \sum x_i} \cdot \mathbf{1}_{\min(x_i) \geq 0} \end{aligned}$$

we need to maximize with respect to λ , since $\mathbf{1}_{\min(x_i)}$ does not depend on λ we can forget it, and consider:

$$h(\lambda) = \lambda^n e^{-\lambda \sum x_i}$$

taking log in both sides, with get a better expression to deal with.

$$g(\lambda) = \log(h(\lambda)) = \log(\lambda^n e^{-\lambda \sum x_i}) = n \log(\lambda) - \lambda \sum x_i$$

Since

$$g'(\lambda) = \frac{n}{\lambda} - \sum x_i \text{ then } \lambda = \frac{n}{\sum x_i} = \frac{1}{\bar{X}_n} \text{ is critical point.}$$

and $g''(\lambda) = \frac{-n}{\lambda^2} < 0$ then this critical point correspond to a maximum.

So, $\lambda = \frac{1}{\bar{X}_n}$ is solution of the maximization problem and therefore the *Maximum likelihood estimator* is:

$$\hat{\lambda}_n = \frac{n}{\sum X_i}$$

2. Let consider X_1, \dots, X_n *i.i.d* $\mathcal{U}([0, \theta])$

- Compute the maximum likelihood estimator.
- Compute the bias of this estimator and transform it into an unbiased one.

Solution:

Compute the maximum likelihood estimator:

$$\begin{aligned} \mathcal{L}(x_1, \dots, x_n, \theta) &= \prod_{i=1}^n f_{\theta}(x_i) \\ &= \prod_{i=1}^n \frac{1}{\theta} \cdot \mathbf{1}_{[0, \theta]}(x_i) \\ &= \frac{1}{\theta^n} \cdot \mathbf{1}_{\min(x_i) \geq 0} \cdot \mathbf{1}_{\max(x_i) \leq \theta} \end{aligned}$$

Since $\mathbf{1}_{\min(x_i) \geq 0}$ does not depend on λ we can forget it, and consider:

$$g(\lambda) = \frac{1}{\theta^n} \cdot \mathbf{1}_{\max(x_i) \leq \theta} = \frac{1}{\theta^n} \cdot \mathbf{1}_{[\max(x_i), +\infty[}(\theta)$$

we can't compute g' since this function it's not derivable in all points.

But since the function $\frac{1}{t^n}$ it's decreasing we can conclude that the maximum estimator likelihood is given by:

$$\hat{\theta}_n = \max(X_i)$$

Compute the bias of this estimator and transform it into an unbiased one.

In order to compute the bias of this estimator, we have to compute: $\mathbb{E}(\hat{\theta}_n) = \mathbb{E}(\max(X_i))$.

To do this we need to determine the density function of this new random variable $\max(X_i)$.

How to compute a density?

1. First step: Computation of the distribution function.

Let $t \in \mathbb{R}$.

$$\begin{aligned} F_{\hat{\theta}_n}(t) &= P(\hat{\theta}_n \leq t) \\ &= P(\max(X_i) \leq t) \\ &\iff \\ &P(X_1 \leq t, \dots, X_n \leq t) \\ &= \prod_{i=1}^n P(X_i \leq t) \quad (\text{by independency of the r.v}) \\ &= (P(X_1 \leq t))^n \quad (\text{because they are identically distributed.}) \end{aligned}$$

Since

$$P(X_1 \leq t) = \int_{-\infty}^t \frac{1}{\theta} \mathbb{1}_{[0, \theta]}(x) dx = \begin{cases} 0 & t < 0 \\ \frac{t}{\theta} & t \in [0, \theta] \\ 1 & t > \theta \end{cases}$$

We conclude:

$$F_{\hat{\theta}_n}(t) = \begin{cases} 0 & t < 0 \\ \frac{t^n}{\theta^n} & t \in [0, \theta] \\ 1 & t > \theta \end{cases}$$

2. Second step: Computation of the density function.

Using the relation between the distribution and density function we have:

$$f_{\hat{\theta}_n}(t) = F'_{\hat{\theta}_n}(t) = n \cdot \frac{t^{n-1}}{\theta^n} \cdot \mathbf{1}_{[0, \theta]}(t)$$

3. Third step: Computation of $\mathbb{E}(\hat{\theta}_n)$

$$\mathbb{E}[\hat{\theta}_n] = \int_{\mathbb{R}} x \cdot f_{\hat{\theta}_n}(x) dx = \frac{n}{n+1} \theta$$

Now we have:

$$b(\hat{\theta}_n) = \mathbb{E}[\hat{\theta}_n] - \theta = -\frac{\theta}{n+1} \neq 0$$

Therefore $\hat{\theta}_n$ is not an unbiased estimator for θ . But since $\lim_{n \rightarrow \infty} b(\hat{\theta}_n) = 0$ it is an asymptotically unbiased estimator for θ .

How to obtain an unbiased estimator?

- First attempt:

We have $\mathbb{E}[\hat{\theta}_n] - \theta = -\frac{\theta}{n+1}$, then:

$$\begin{aligned} \mathbb{E}[\hat{\theta}_n] - \theta + \frac{\theta}{n+1} &= 0 \\ \Leftrightarrow \\ \mathbb{E}\left[\hat{\theta}_n + \frac{\theta}{n+1}\right] - \theta &= 0 \end{aligned}$$

It is ok to consider $\tilde{\theta}_n = \hat{\theta}_n + \frac{\theta}{n+1}$ as unbiased estimator of θ ?

The answer is **no**, $\tilde{\theta}_n$ depends on an unknown parameter (θ in this case), so by definition it is not an estimator!

- Second attempt:

We have:

$$\begin{aligned}
\mathbb{E}[\hat{\theta}_n] - \theta &= -\frac{\theta}{n+1} = 0 \\
&\iff \\
\mathbb{E}[\hat{\theta}_n] &= \theta - \frac{\theta}{n+1} \\
&\iff \\
\mathbb{E}[\hat{\theta}_n] &= \frac{n}{n+1}\theta \\
&\iff \\
\mathbb{E}\left[\frac{n+1}{n}\hat{\theta}_n\right] &= \theta \\
&\iff \\
\mathbb{E}\left[\frac{n+1}{n}\hat{\theta}_n\right] - \theta &= 0
\end{aligned}$$

Thus, $\frac{n+1}{n}\hat{\theta}_n$ is an unbiased estimator for θ .

```
## Comparing the estimators for the uniform distribution.
```

```
n= 100
```

```
U = runif(n, 0, 4)
```


```
theta = max(U)
```

```
for (i in 1:50) {
  U = runif(n, 0, 4)
  theta = c(theta, max(U))
}
```

```
mean(theta)
```

```
## [1] 3.963466
```

```
boxplot(theta)
```




```
## Adjust to make the estimator unbiased
```

```
thetab = (n+1)/n*theta
```

```
mean(thetab)
```

```
## [1] 4.003101
```

```
boxplot(thetab)
```



```
## With more observations
```

```
n= 100
```

```
U = runif(n, 0, 4)
```


```
theta = max(U)
```

```
for (i in 1:5000) {
  U = runif(n, 0, 4)
  theta = c(theta, max(U))
}
```

```
mean(theta)
```

```
## [1] 3.960618
```

```
boxplot(theta)
```



Property Let X_1, \dots, X_n *i.i.d* random variables.

- Let $\mu = \mathbb{E}[X_1]$ (unknown)
- Let $\sigma^2 = V[X_1]$ (unknown)

A classical estimator for:

- μ is

$$\hat{\mu}_n = \bar{X}_n$$

- σ^2 is

$$\hat{\sigma}_n^2 = \frac{1}{n} \sum (X_i - \bar{X}_n)^2$$

Exercise

Show that:

1. $\hat{\mu}$ is unbiased.
2. $\hat{\sigma}_n^2$ is biased and that $\frac{n}{n-1} \hat{\sigma}_n^2$ is unbiased.

Solution:

1.

$$\begin{aligned}\mathbb{E}[\hat{\mu}_n] &= \mathbb{E}\left[\frac{1}{n} \sum X_i\right] \\ &= \frac{1}{n} \sum \mathbb{E}[X_i] \quad (\text{by linearity}) \\ &= \frac{1}{n} n \cdot \mathbb{E}[X_1] \quad (\text{as they are indentially distributed}) \\ &= \mathbb{E}[X_1]\end{aligned}$$

Therefore: $\hat{\mu}_n = \bar{X}_n$ is an unbiased estimator for μ .

2.

First note that we can rewrite $\hat{\sigma}_n^2$ as follows:

$$\begin{aligned}\hat{\sigma}_n^2 &= \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X}_n)^2 \\ &= \frac{1}{n} \sum_{i=1}^n (X_i^2 - 2X_i\bar{X}_n + \bar{X}_n^2) \\ &= \frac{1}{n} \sum_{i=1}^n X_i^2 - 2\bar{X}_n \left(\frac{1}{n} \sum_{i=1}^n X_i \right) + \bar{X}_n^2 \quad (\text{check that factors that depends on } i) \\ &= \frac{1}{n} \sum_{i=1}^n X_i^2 - 2\bar{X}_n^2 + \bar{X}_n^2 \\ &= \frac{1}{n} \sum_{i=1}^n X_i^2 - \bar{X}_n^2\end{aligned}$$

Now use this expression to compute the expectation:

$$\begin{aligned}\mathbb{E}[\hat{\sigma}_n^2] &= \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n X_i^2 - \bar{X}_n^2\right] \\ &= \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i^2] - \mathbb{E}[\bar{X}_n^2] \quad (\text{by linearity}) \\ &= \frac{1}{n} \sum_{i=1}^n (V[X_i] + \mathbb{E}[X_i]^2) - (V[\bar{X}_n] + \mathbb{E}[\bar{X}_n]^2) \quad (\text{by using the variance formula}) \\ &= \sigma^2 + \mu^2 - V[\bar{X}_n] - \mu^2 \quad (\text{by substituting the expressions of the variance and expectation}) \\ &= \sigma^2 - V[\bar{X}_n]\end{aligned}$$

Now lets compute $V[\bar{X}_n]$:

$$\begin{aligned}
V[\bar{X}_n] &= V\left[\frac{1}{n} \sum_{i=1}^n X_i\right] \\
&= \frac{1}{n^2} V\left[\sum_{i=1}^n X_i\right] \quad (\text{by linearity}) \\
&= \frac{1}{n^2} \sum_{i=1}^n V[X_i] \quad (\text{because they are independent}) \\
&= \frac{\sigma^2}{n}
\end{aligned}$$

Putting all together we get:

$$b(\hat{\sigma}_n^2) = \mathbb{E}[\hat{\sigma}_n^2] - \sigma^2 = -V[\bar{X}_n] = -\frac{\sigma^2}{n}$$

so $\hat{\sigma}_n^2$ is an biased estimator.

But we can modify it to obtain an unbiased version by noticed that:

$$\mathbb{E}\left[\frac{n}{n-1} \hat{\sigma}_n^2\right] = \sigma^2$$

Therefore:

$$\hat{\sigma}_{n-1}^2 = \frac{n}{n-1} \hat{\sigma}_n^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X}_n)^2$$

is an unbiased estimator for σ^2 .

Theorem Let X_1, \dots, X_n *i.i.d* with $\mu = \mathbb{E}[X_1]$ and $\sigma^2 = V[X_1]$.

- An unbiased estimator for μ is:

$$\hat{\mu}_n = \bar{X}_n \quad \text{and its variance is} \quad \frac{\sigma^2}{n}$$

- An unbiased estimator for σ^2 is:

$$\hat{\sigma}_{n-1}^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X}_n)^2$$

Quality of an estimator

Mean Quadratic Error Definition: Let θ an unknown parameter, let $\hat{\theta}_n$ an estimator of θ , the mean quadratic error (MQE) is given by:

$$MQE(\hat{\theta}_n) = \mathbb{E}[(\hat{\theta}_n - \sigma)^2]$$

Property:

$$MQE = V[\hat{\theta}_n] + ((b(\hat{\theta}_n)))^2$$

Proof:

Lets rewrite the MQE definition.

$$\begin{aligned} MQE(\hat{\theta}_n) &= \mathbb{E}[(\hat{\theta}_n - \sigma)^2] \\ &= \mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n] + \mathbb{E}[\hat{\theta}_n] - \sigma)^2] \quad (\text{by adding and subtracting yhe same thing}) \\ &= \underbrace{\mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2]}_{V[\hat{\theta}_n]} + \underbrace{\mathbb{E}[(\mathbb{E}[\hat{\theta}_n] - \sigma)^2]}_{\mathbb{E}[b(\hat{\theta}_n)]^2} + 2 \underbrace{\mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])(\mathbb{E}[\hat{\theta}_n] - \sigma)]}_{A_n} \end{aligned}$$

So to proof our result we just need to proof that $A_n = 0$.

$$\begin{aligned} A_n &= \mathbb{E}[\underbrace{(\mathbb{E}[\hat{\theta}_n] - \sigma)}_{b(\hat{\theta}_n)} \cdot (\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])] \\ &= \underbrace{b(\hat{\theta}_n)}_{\text{deterministic term}} \cdot \mathbb{E}[\underbrace{(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])}_{\text{this is a centeret r.v.}}] \quad (\text{the expectation of a centered r.v is zero.}) \\ &= 0 \end{aligned}$$

Comparison of estimators Definition: Let θ an unknown parameter, $\hat{\theta}_{n,1}$ and $\hat{\theta}_{n,2}$ for θ .

$\hat{\theta}_{n,1}$ is better than $\hat{\theta}_{n,2}$ if

$$\forall n \geq 1, \quad MQE(\hat{\theta}_{n,1}) \leq MQE(\hat{\theta}_{n,2})$$

Application Let X_1, \dots, X_n *i.i.d* $\mathcal{U}([0, \theta])$, compare the following estimators to see if one is better than the others.

- $\hat{\theta}_{n,1} = 2\bar{X}_n$
- $\hat{\theta}_{n,2} = \max(X_n)$
- $\hat{\theta}_{n,3} = \frac{n+1}{n}\bar{X}_n$

Remember:

$$f_{\hat{\theta}_{n,2}}(t) = n \frac{t^{n-1}}{\theta^n} \mathbf{1}_{[0, \theta]}(t)$$

Solution:

Lets compute the MQE of each estimator.

1. Since $\hat{\theta}_{n,1}$ is an unbiased estimator for θ , we have:

$$\begin{aligned}
MQE(\hat{\theta}_{n,1}) &= V[\hat{\theta}_{n,1}] \\
&= V[2\bar{X}_n] \\
&= 4V[\bar{X}_n] \\
&= 4 \sum_{n=1}^n V[X_1] \quad (\text{by using i.i.d}) \\
&= 4 \frac{\sigma^2}{12n} \quad (\text{check the variance of a uniform distribution})
\end{aligned}$$

Thus:

$$MQE(\hat{\theta}_{n,1}) = \frac{\sigma^2}{3n}$$

2. We already know that: $b(\hat{\theta}_{n,2}) = -\frac{\theta}{n+1}$ and $\mathbb{E}[\hat{\theta}_{n,2}] = \frac{n}{n+1}\theta$.

We have to compute $V[\hat{\theta}_{n,2}]$, since

$$V[\hat{\theta}_{n,2}] = \mathbb{E}[(\hat{\theta}_{n,2})^2] - \mathbb{E}[\hat{\theta}_{n,2}]^2$$

and

$$\mathbb{E}[(\hat{\theta}_{n,2})^2] = \int_{\mathbb{R}} n \frac{t^{n-1}}{\theta^n} \mathbf{1}_{[0,\theta]}(t) dt = \frac{n}{n+2} \theta^2$$

Therefore:

$$\begin{aligned}
MQE(\hat{\theta}_{n,2}) &= V[\hat{\theta}_{n,2}] - b(\hat{\theta}_{n,2})^2 \\
&= \frac{n}{n+2} \theta^2 - \frac{n^2}{(n+1)^2} \theta^2 + \frac{1}{(n+1)^2} \theta^2 \\
&= \frac{2}{(n+1)(n+2)} \theta^2
\end{aligned}$$

3. Since $\hat{\theta}_{n,3}$ is an unbiased estimator for θ , we have:

$$\begin{aligned}
MQE(\hat{\theta}_{n,3}) &= V[\hat{\theta}_{n,3}] \\
&= V\left[\frac{n+1}{n} \hat{\theta}_{n,2}\right] \\
&= \frac{(n+1)^2}{n^2} V[\hat{\theta}_{n,2}] \\
&= \frac{(n+1)^2}{n^2} \cdot \frac{n}{(n+1)^2(n+2)} \theta^2 \\
&= \frac{1}{n(n+2)} \theta^2
\end{aligned}$$

It is easy to check that for $n \geq 1$:

$$3n \leq \frac{(n+1)}{2}(n+2) \leq n(n+2)$$

Thus

$$MQE(\hat{\theta}_{n,3}) \leq MQE(\hat{\theta}_{n,2}) \leq MQE(\hat{\theta}_{n,1})$$

and we conclude that $MQE(\hat{\theta}_{n,3})$ it's the best estimator for θ .

Theorem Let θ an unknown estimator and $\hat{\theta}_n$ an unbiased estimator of θ .

Cramer-Rao bound:

$$V[\hat{\theta}_n] \geq \frac{1}{I_n(\theta)}$$

where

$$I_n(\theta) = -\mathbb{E}\left[\frac{\partial^2 \log(\mathcal{L}(x_1, \dots, x_n; \theta))}{\partial \theta^2}\right]$$

is the *Fisher Information*.

This is valid under certain conditions:

- $I_n(\theta)$ should exist.
- The possible values of X_i should not depend on θ . ($\mathcal{U}([0, \theta])$ not works.)

Improve estimators.

1. If you just have one estimator $\hat{\theta}_n$, you consider:

$$\lambda \cdot \hat{\theta}_n$$

with $\lambda \in \mathbb{R}^*$ and find λ that minimize:

$$MQE(\lambda \hat{\theta}_n)$$

2. If you have two estimators $\hat{\theta}_{n,1}, \hat{\theta}_{n,2}$.

- Construct:

$$\hat{\theta}_{n,a,b} = a\hat{\theta}_{n,1} + b\hat{\theta}_{n,2}$$

- Compute:

$$M_{ab} := MQE(\hat{\theta}_{n,a,b})$$

- Minimize M_{ab} in (a, b) .

If $\hat{\theta}_{n,1}$ and $\hat{\theta}_{n,2}$ are **unbiased** estimators and you want to keep $\hat{\theta}_{n,a,b}$ unbiased, take $b = 1 - a$.

Exercise Let X_1, \dots, X_n i.i.d $\mathcal{P}(\lambda)$.

Compute an estimator for λ by:

1. Method of moments.
2. Maximum likelihood.

Solution

1. Since $X \sim \mathcal{P}(\lambda) \Rightarrow \begin{cases} \mathbb{E}[X] = \lambda \\ V[X] = \lambda \end{cases}$

One estimator is:

$$\hat{\lambda}_{n,1} = \mathbb{E}[X] = \frac{1}{n} \sum_{i=1}^n X_i$$

another estimator for λ is given by:

$$\hat{\lambda}_{n,2} = \mathbb{E}[X] = \frac{1}{n} \sum_{i=1}^n (X_i - \hat{X}_n)^2$$

2. Let x_1, \dots, x_n positive integers.

$$\begin{aligned} \mathcal{L}(x_1, \dots, x_n; \lambda) &= \prod_{i=1}^n P(X_1 = x_i) \quad (\text{identically distributed}) \\ &= \prod_{i=1}^n \left(\frac{\lambda^{x_i} e^{-\lambda}}{x_i!} \right) \\ &= \lambda^{\sum x_i} e^{-n\lambda} \frac{1}{\underbrace{\prod_{i=1}^n x_i!}_{\text{Not depends on } \lambda}} \end{aligned}$$

since we need to minimize over λ and log is an increasing function, we consider:

$$g(\lambda) = \sum x_i \cdot \log(\lambda) - n\lambda$$

now, find the critical to find the critical points, we solve

$$g'(\lambda) = 0 \iff \lambda = \frac{1}{n} \sum x_i$$

by analysing the sign of g' we can conclude that this point is a maximum and therefore:

$$\hat{\lambda}_{n,3} = \frac{1}{n} \sum_{i=1}^n X_i$$

is an estimator for λ (the same obtained by the Methods of Moments).

Practical Work

1. Load the data in R software.
2. Propose a model for the variables associated to this file.

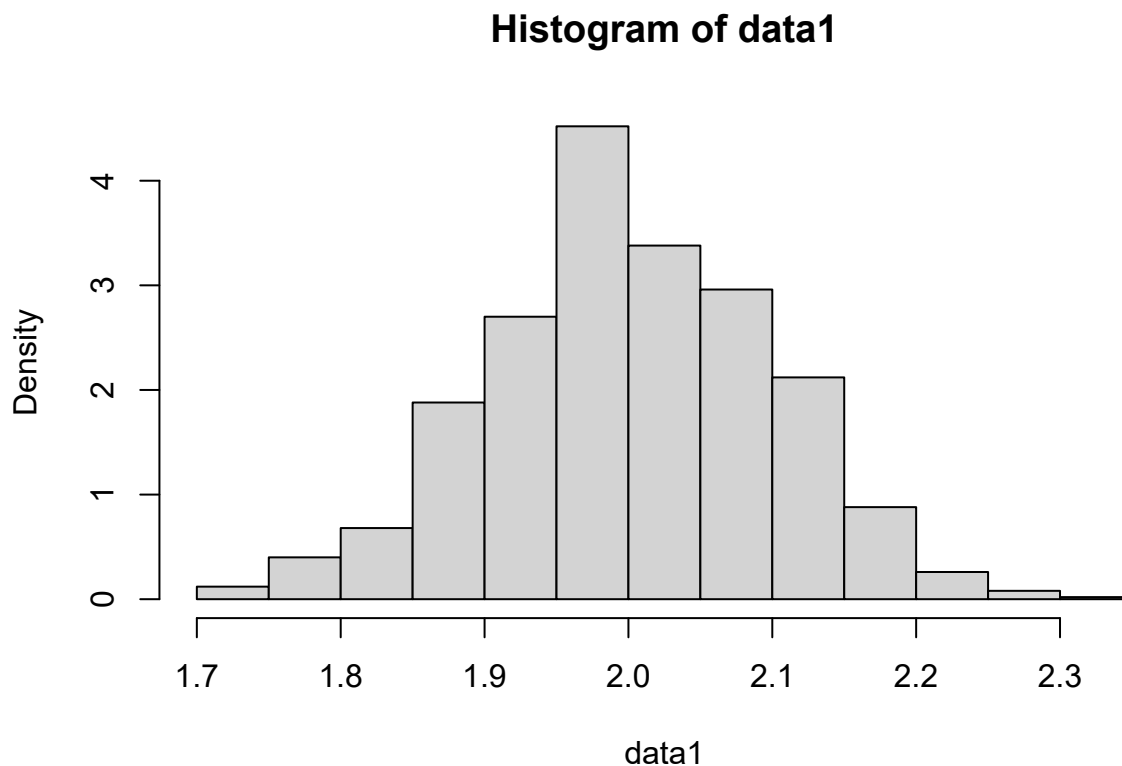
- Make a visualization of this.

Data 1 Load the data.

```
data1 <- as.matrix(read.table("data/data1.txt"))  
  
## Some comments about the data and data types.
```

In this two graph we see that we can't determine just by looking at the graph a precise value for the parameters.

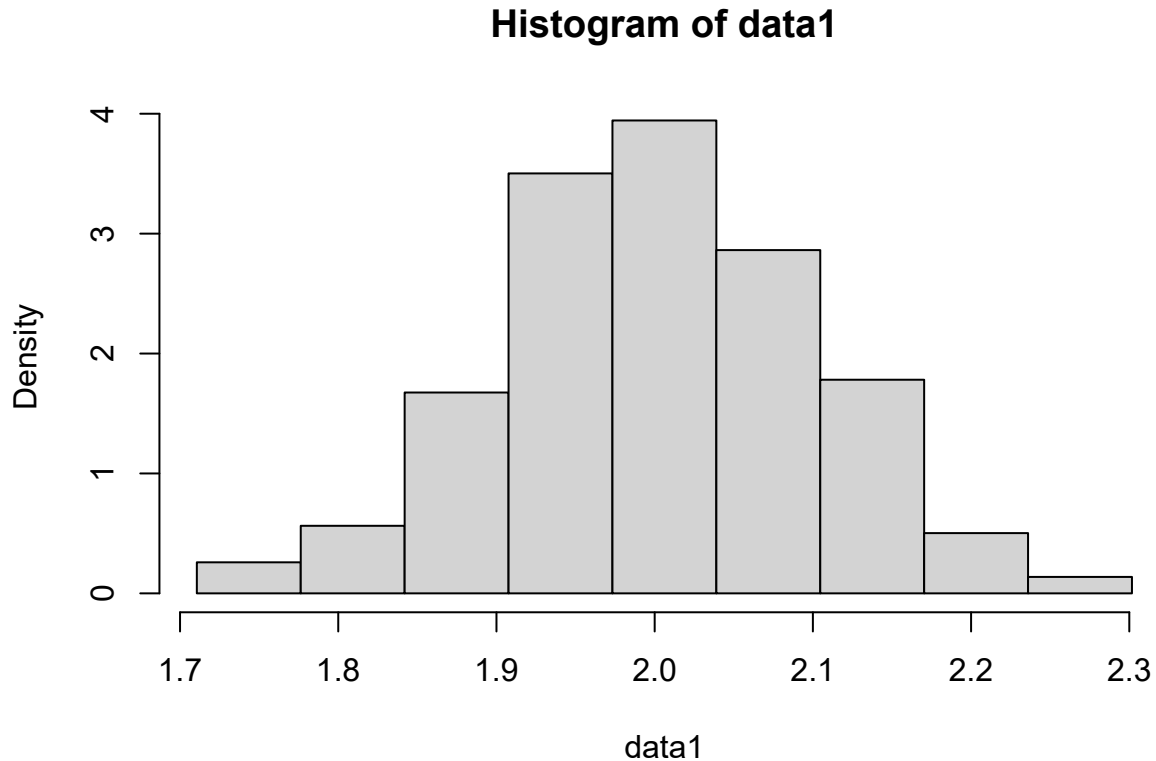
```
## Set freq = FALSE in order to obtain a correct graph.  
  
## Be careful about getting conclusions, since the function make the  
## bins automatically.  
  
hist(data1, freq = FALSE)
```



```
## Create a personalized breaks

data1_min = min(data1)
data1_max = max(data1)
l = seq(data1_min, data1_max, length = 10)

hist(data1, freq = FALSE, breaks = l)
```



We guess a Gaussian distribution.

How to estimate the parameters of the distribution?

1. Compute the empirical mean and statistical variance.
2. Plot the theoretical density that we guess.
3. Test the goodness of fitness.

Compute the empirical mean and statistical variance.

```
mhu <- mean(data1)

# When you compute the variance be careful and read the documentation
# in this case we have:
# "The denominator n - 1 is used which gives an unbiased estimator of the (co)variance
# for i.i.d. observations"
```

```
sigma2 <- var(data1)

## You also can do it manually:

sigma2 <- 1/(nrow(data1)-1)*sum((data1 - mean(data1))^2)
```

Plot the theoretical density that we guess.

To this graph we use the information of the histogram associated to the data.

```
## This generates a list, with all the hist information.

H <- hist(data1, freq = FALSE, plot = FALSE) ## ignore warning message
H

## $breaks
## [1] 1.70 1.75 1.80 1.85 1.90 1.95 2.00 2.05 2.10 2.15 2.20 2.25 2.30 2.35
##
## $counts
## [1] 6 20 34 94 135 226 169 148 106 44 13 4 1
##
## $density
## [1] 0.12 0.40 0.68 1.88 2.70 4.52 3.38 2.96 2.12 0.88 0.26 0.08 0.02
##
## $mids
## [1] 1.725 1.775 1.825 1.875 1.925 1.975 2.025 2.075 2.125 2.175 2.225 2.275
## [13] 2.325
##
## $xname
## [1] "data1"
##
## $equidist
## [1] TRUE
##
## attr(,"class")
## [1] "histogram"
```

Comparison graph

```
## Histogram info to use.
limits <- H$breaks
lmin <- limits[1]
lmax <- limits[length(limits)]

## Create our plot

## Create a vector that represents the [lmin, lmax] interval.
x <- seq(lmin, lmax, by = 0.01)

## Compute the density of the vector x associated to the parameters mhu y sigma2.
y <- dnorm(x, mhu, sqrt(sigma2))

## Set the y_max as the max in the empirical density and the 'artificial' density.
```


```

y_max = max(y, H$density)

## Make a plot for the comparison

hist(data1, freq = FALSE, xlim = c(lmin, lmax), ylim = c(0, y_max*1.01))
par(new = TRUE)
plot(x, y, type = 'l', col = 'red'
     , xlim = c(lmin, lmax)
     , ylim = c(0, y_max*1.01)
     , xlab = ""
     , ylab = "")

```



Test the goodness of fitness.

```
ks.test(x, 'pnorm', mhu, sqrt(sigma2))
```

```

##
## One-sample Kolmogorov-Smirnov test
##
## data: x
## D = 0.25147, p-value = 0.0003644
## alternative hypothesis: two-sided

```

```
## Outputs:
```

```
## The D: notion of distance between the data and the guess distribution.  
## The p-value: Used to reject or accept the propose model.
```

Model values vs Real values

```
## Model values  
#mhu  
#sigma2
```

```
## Real values  
real_mhu = 2.0  
real_sigma = 0.1
```

```
cat('model_mhu:', mhu, 'vs', 'real_mhu:', real_mhu, '\n')
```

```
## model_mhu: 2.001483 vs real_mhu: 2
```

```
cat('model_sigma:', sqrt(sigma2), 'vs', 'real_sigma:', real_sigma)
```

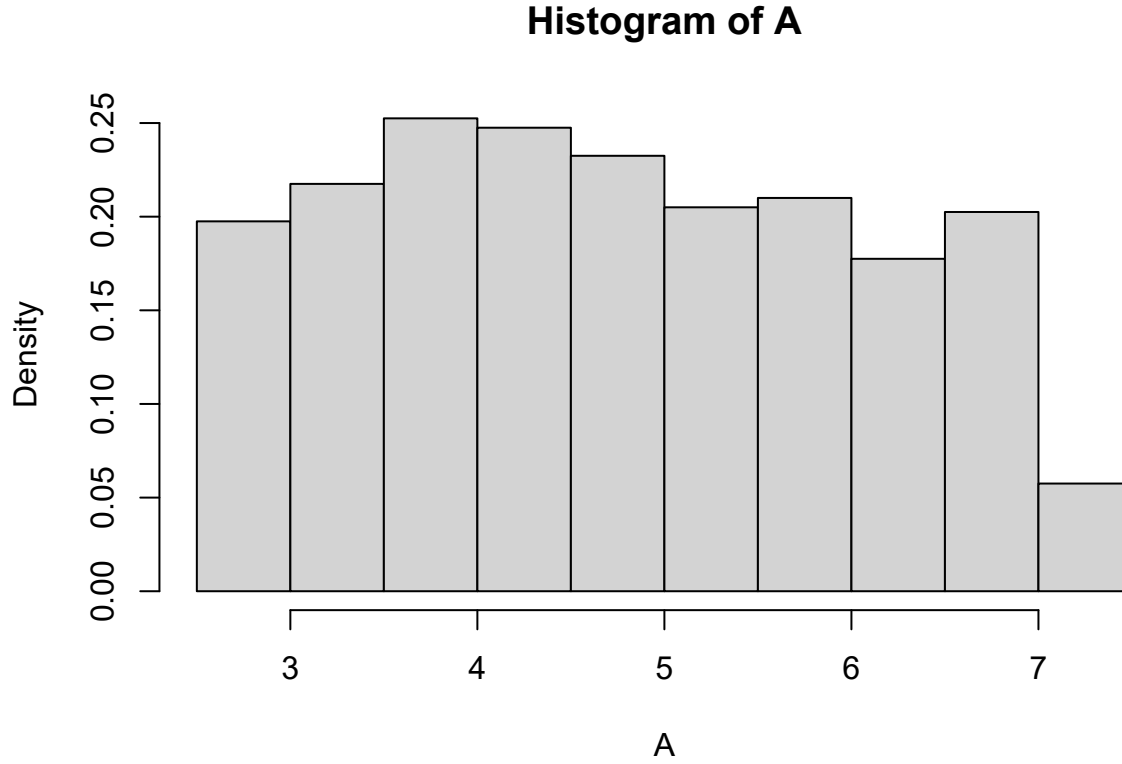
```
## model_sigma: 0.09946156 vs real_sigma: 0.1
```

Data 2 Load the data.

```
## Loading the data  
A <- as.matrix(read.table("data/data2.txt"))  
  
A_max = max(A)
```

In this graph we can observe that there is potentially less observation in the last class, but if we look at the maximum of our dataset (7.099357), this last class exist because the maximum is slighter bigger than 7, so R automatically creates a new class just for this element.

```
## Make visualizations  
hist(A, freq = FALSE)
```

To avoid this, we can use a strategy to create better classes.

Create classes If we have a dataset with n observations.

The number of classes is given by:

$$K \approx 1 + 3.22 \times \log_{10}(n) \quad (\text{take the floor to get an integer value.})$$

The limits of the classes are given by:

$$a_0 = \min(x_i) - \varepsilon \quad a_K = \max(x_i) + \varepsilon$$

the ε correction is to ensure that all the observations are contained in our partition.

Also the ε value depends on the order of magnitude of the data (numerator of the equation) and the precision of the software that we are using (denominator of the equation).

$$\varepsilon = \frac{\max(x_i) - \min(x_i)}{10^9}$$

the other break points can be constructed as follows:

$$a_i = a_0 + i \times \text{amp} \quad \text{for } i \in \{1, \dots, K\}$$

where

$$amp = \frac{a_K - a_0}{K} \quad (\text{length of the classes})$$

Lets use this in our dataset.


```
## Compute correct classes parameters:
n = length(A)
K = floor(1 + 3.22*log(n, 10))
A_min = min(A)
A_max = max(A)

amp = (A_max - A_min)/K

epsilon = (A_max - A_min)/10^9

limits = seq(A_min, A_max, by = amp)

hist(A, freq = FALSE, breaks = limits)
```



We guess is an uniform distribution $\mathcal{U}[a, b]$.

In this cases we have two estimate 2 parameters, previously we had been working with the uniform distribution $\mathcal{U}[0, \theta]$ and prove that $2\bar{X}_n$ and $\frac{n+1}{n} \max(X_i)$ are unbiased and $\max(X_i)$ is a biased estimator for θ .

In order to find unbiased estimators for the parameters a, b of our distribution $\mathcal{U}[a, b]$ we should make all the computations again.

In practice we can consider the *naive* estimators:

$$a = \min(X_i) \quad b = \max(X_i)$$

or we can try to find unbiased estimators using the Method of Moments:

Let $X \sim \mathcal{U}[a, b]$, so:

$$\mathbb{E}[X] = \frac{a+b}{2} \quad \text{and} \quad V[X] = \frac{(b-a)^2}{12}$$

By applying the method of moments we get:

$$\frac{\hat{a}_n + \hat{b}_n}{2} = \bar{X}_n \quad \text{and} \quad \frac{(\hat{b}_n - \hat{a}_n)^2}{12} = \hat{\sigma}_n^2$$

solving this system for \hat{a}_n and \hat{b}_n we obtain:

$$\hat{a}_n = \frac{1}{2} \left(2\bar{X}_n - \sqrt{12\hat{\sigma}_n^2} \right) \quad \hat{b}_n = \frac{1}{2} \left(2\bar{X}_n + \sqrt{12\hat{\sigma}_n^2} \right)$$

For this case we are going to use the Naive estimators, and compare again the Model values vs Real values.

```
## Model values
a = min(A)
b = max(A)

## Real values
real_a = 2.5
real_b = 7.1

cat('model_a:', a, 'vs', 'real_a:', real_a, '\n')
```

```
## model_a: 2.503019 vs real_a: 2.5
```

```
cat('model_b:', b, 'vs', 'real_b:', real_b)
```

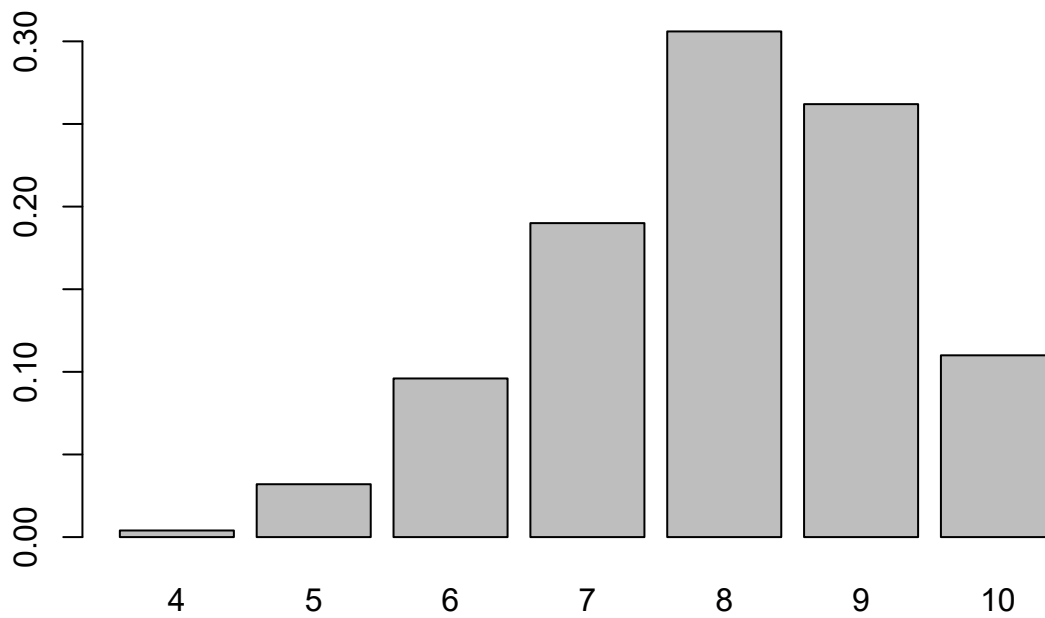
```
## model_b: 7.099357 vs real_b: 7.1
```

Data 3 In progress

Data 4 In progress

```
## Loading the data
A <- as.matrix(read.table("data/data4.txt"))

barplot(table(A)/length(A))
```



```
## A binomial distribution can be approach by a Gaussian distribution.
```

```
## Methods of moments
```

```
## Don't forget to correct the estimation of np to be an integer.
```

```
## and then correct the ph value
```

```
m = mean(A)
s2 = var(A)*(499/500)
```

```
ph = 1-s2/m
```

```
nh = m/ph
```

```
nh
```

```
##          V1
```

```
## V1 10.0325
```

```
nh = 10
```

```
ph = m/nh
```

```
ph
```


```
## [1] 0.7988
```

Data 5 In progress

```
## Loading the data
```

```
A <- as.matrix(read.table("data/data5.txt"))
```

```
barplot(table(A)/length(A))
```



```
## A binomial distribution can be approach by a Gaussian distribution.
```

```
## Methods of moments
```

```
## Don't forget to correct the estimation of np to be an integer.
```

```
## and then correct the ph value
```

```
m = mean(A)
```

```
s2 = var(A)*(499/500)
```

```
ph = 1-s2/m
```

```
nh = m/ph
```

```
nh
```

```
## V1
```

```
## V1 7.38796
```

```
nh = 7
```

```
ph = m/nh
```

```
ph
```

```
## [1] 0.1342857
```

```
## When p is small is very hard to get correct estimations of n when you  
## don't have enough number of observations
```

Day 3 (In progress)


Confidence interval

```
m1 = c()
```

```
for (i in 1:50) {  
  A = rnorm(500, 2, 1)
```

```
  m1 = c(m1, mean(A))  
}
```

```
boxplot(A)
```




```
## Increase the number of observations
```


```
m2 = c()
```

```
for (i in 1:50) {  
  B = rnorm(2000, 2, 1)  
  
  m2 = c(m2, mean(B))  
}
```

```
boxplot(B)
```



```
## Change the variance  
m3 = c()  
  
for (i in 1:50) {  
  C = rnorm(2000, 2, 10)  
  
  m3 = c(m3, mean(C))  
}  
  
boxplot(C)
```

```
## Create a data frame with all the observations
```


```
a = c(m1, m2, m3)
```

```
r = rep(c("A", "B", "C"), each = 50)
```

```
S = data.frame(x = a, l = r)
```

In this plot we can see that the approximation depends on the number of observations and the variance of the sample.

```
boxplot(S$x ~ S$l)
```




Chi squared distribution

```
x = seq(0.1, 10, by = 0.01)

y1 = dchisq(x, 1)
y2 = dchisq(x, 2)
y3 = dchisq(x, 3)
y4 = dchisq(x, 4)
y6 = dchisq(x, 6)
y9 = dchisq(x, 9)

y_max = max(y1, y2, y3, y4, y6, y9)

plot(x, y1, type = 'l', col = 'yellow', xlim = c(0, 10), ylim = c(0, 0.5))
par(new = TRUE)
plot(x, y2, type = 'l', col = 'green', xlim = c(0, 10), ylim = c(0, 0.5))
par(new = TRUE)
plot(x, y3, type = 'l', col = 'black', xlim = c(0, 10), ylim = c(0, 0.5))
par(new = TRUE)
plot(x, y4, type = 'l', col = 'blue', xlim = c(0, 10), ylim = c(0, 0.5))
par(new = TRUE)
plot(x, y6, type = 'l', col = 'pink', xlim = c(0, 10), ylim = c(0, 0.5))
par(new = TRUE)
plot(x, y9, type = 'l', col = 'red', xlim = c(0, 10), ylim = c(0, 0.5))
```



t-student distribution

```
x = seq(-10, 10, by = 0.01)

y1 = dt(x, 1)
y2 = dt(x, 2)
y3 = dt(x, 3)
y4 = dt(x, 4)
y6 = dt(x, 6)
y20 = dt(x, 20)

y_norm = dnorm(x)


y_max = max(y1, y2, y3, y4, y6, y20)

plot(x, y1, type = 'l', col = 'yellow', xlim = c(-10, 10), ylim = c(0, y_max))
par(new = TRUE)
plot(x, y2, type = 'l', col = 'green', xlim = c(-10, 10), ylim = c(0, y_max))
par(new = TRUE)
plot(x, y3, type = 'l', col = 'black', xlim = c(-10, 10), ylim = c(0, y_max))
par(new = TRUE)
plot(x, y4, type = 'l', col = 'blue', xlim = c(-10, 10), ylim = c(0, y_max))
par(new = TRUE)
plot(x, y6, type = 'l', col = 'pink', xlim = c(-10, 10), ylim = c(0, y_max))
```

```

par(new = TRUE)
plot(x, y20, type = 'l', col = 'red', xlim = c(-10, 10), ylim = c(0, y_max))
par(new = TRUE)
plot(x, y_norm, type = 'l', col = 'black', xlim = c(-10, 10), ylim = c(0, y_max))

```



```
## A t-student distribution with a big enough k approach a Gaussian distribution
```

Day 4 (In progress)

Use of t-student table

Confidence interval

```

## We compute N observations of  $\overline{X}_n$ .

## Number of r.v.
n <- 5

## Number of observations
N <- 50000

## Parameter exp distribution.

```

```

lambda <- 0.7


## Create a matrix
A <- matrix(rexp(n*N, lambda), ncol = n)

## Compute the mean for each row
M <- apply(A, 1, mean)

## Visualization for M
hist(M, freq = F, breaks = 20)

```

Histogram of M



```

###

Mo <- sort(M)
y <- 1:N/N


max_Mo <- max(Mo)

x <- seq(0, max_Mo, by = 0.01)
mu <- 1/lambda
v <- 1/(lambda^2)/n

yt <- pnorm(x, mu, sqrt(v))

```

```
plot(Mo, y, type = 'l', xlim = c(0, max_Mo), ylim = c(0,1))
par(new = TRUE)
plot(x, yt, type = 'l', xlim = c(0, max_Mo), ylim = c(0,1), col = 'red')
```



Appendix

Types of convergence